

Bitcoin

CS 161 Fall 2022 - Lecture 11

Today: Bitcoin

- Bitcoin: Sending and receiving money without trusting a central authority
- How does Bitcoin work?
 - How do we manage identity, transactions, and balances on a ledger?
 - How do we construct a decentralized, trusted ledger?
- The trouble with Bitcoin
- Bitcoin in practice

Problem Statement: The Decentralized Bank

- Alice, Bob, Carol, and Dave each have a sum of currency
- Anyone can send money to anyone else
 - Dave pays Alice 10 coins
 - Dave's balance decreases by 10
 - Alice's balance increases by 10
- No party can spend more currency than they currently have
 - Dave has 10 coins
 - Dave can send Alice 10 coins
 - Dave cannot send Alice 15 coins
- No party trusts any other party, and there is no central authority
 - Usually, a centralized authority (e.g. a bank) tracks balances and enforces spending rules
 - Without a central authority, we must use cryptography to enforce correctness!

Identity Management

- Idea: Use certificates to verify real-world identity?
 - But this needs central root CAs, which we don't want
- Instead, just define each identity as a public key
 - Circumvents the identity problem since we completely ignore real-world identities
 - Instead of transactions between “people,” we use transactions between public keys

Alice



PK_A

Bob



PK_B

Carol



PK_C

Dave



PK_D

Transactions

- *For now*, assume the existence of a trusted, public ledger
 - Accessible to all parties: Everyone can view the ledger
 - Append-only: Everyone can add data to the ledger
 - Immutable: Nobody can change or delete existing data
 - This is a central authority for now, but we'll deal with it later
- How do we record transactions in a way that everyone can see them?
 - Idea: Put " PK_D paid PK_A 10 coins" on the ledger?
 - Problem: Mallory can forge a transaction (e.g. " PK_A paid PK_M 10,000 coins")
 - Solution: Use digital signatures: Put $\{\text{"}PK_D \text{ paid } PK_A \text{ 10 coins"}\}_{SK_D^{-1}}$ on the ledger
 - Problem: How do we check how much currency Dave (PK_D) has to spend?

Transactions

- Each party's balance exists only as (total received) - (total spent) on the ledger
 - Real-time balances values aren't listed directly on the ledger!
- How much time does this take?
 - Too long: We have to scan through the entire ledger to determine someone's balance
 - The ledger grows indefinitely as more and more transactions occur

The Ledger v1

1. PK_A, PK_B, PK_C , and PK_D magically start with 10 coins.
2. $\{\text{"}PK_A \text{ paid } PK_C \text{ 4 coins.}\} SK_A^{-1}$
3. $\{\text{"}PK_B \text{ paid } PK_D \text{ 6 coins.}\} SK_B^{-1}$
4. $\{\text{"}PK_B \text{ paid } PK_A \text{ 2 coins.}\} SK_B^{-1}$

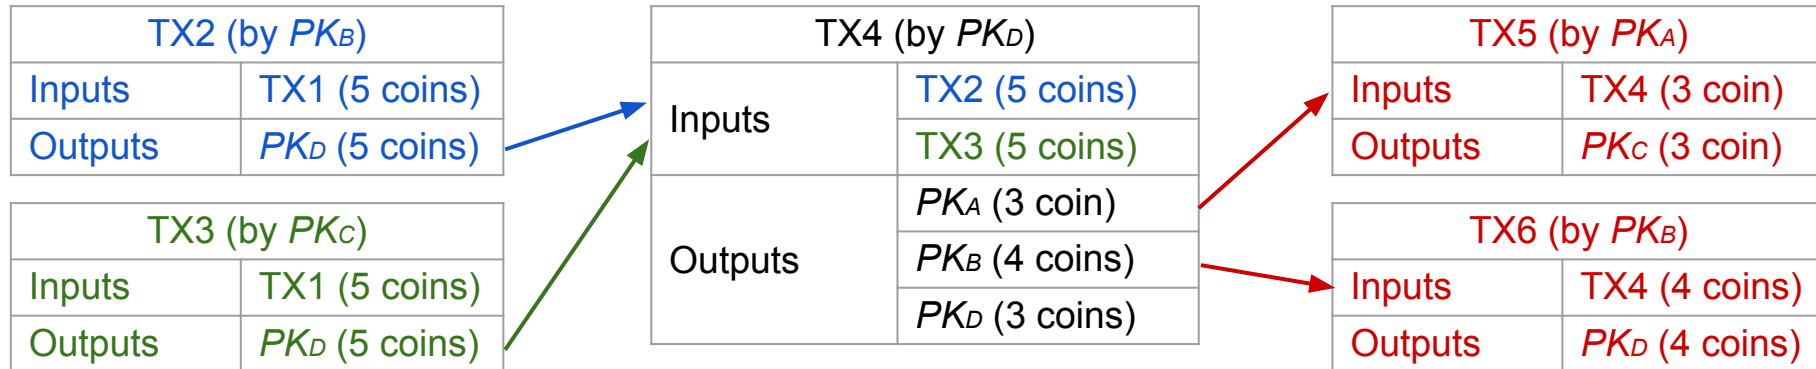
How many coins does PK_A have right now?

8 coins! $10 - 4 + 2 = 8$, for transactions 1, 2, and 4 on the ledger.

Transactions

- Construction: Each transaction has **inputs** (sources, where the money came from) and **outputs** (destinations, who currency is going to)
 - Now, each party only needs to keep track of transactions where they received money (they were the output)
 - Example: {"Using currency from TX 2 and TX 4, give 3 coins to PK_A and 4 coins to PK_B "} $_{SK_D^{-1}}$
 - TX 2 and TX 4 are now "spent" (used as an input), so we don't need to keep track of it anymore
 - Alice and Bob now have an additional unspent transaction
 - To validate a transaction T is unspent, check that no transaction after T uses T as an input
 - Pay whatever change you have back to yourself
 - Example: From TX 2 and TX 4, Dave has received 10 coins in total
 - {"Using TX 2 and TX 4, give 3 coins to PK_A , 4 coins to PK_B , and 3 coins to PK_D "} $_{SK_D^{-1}}$

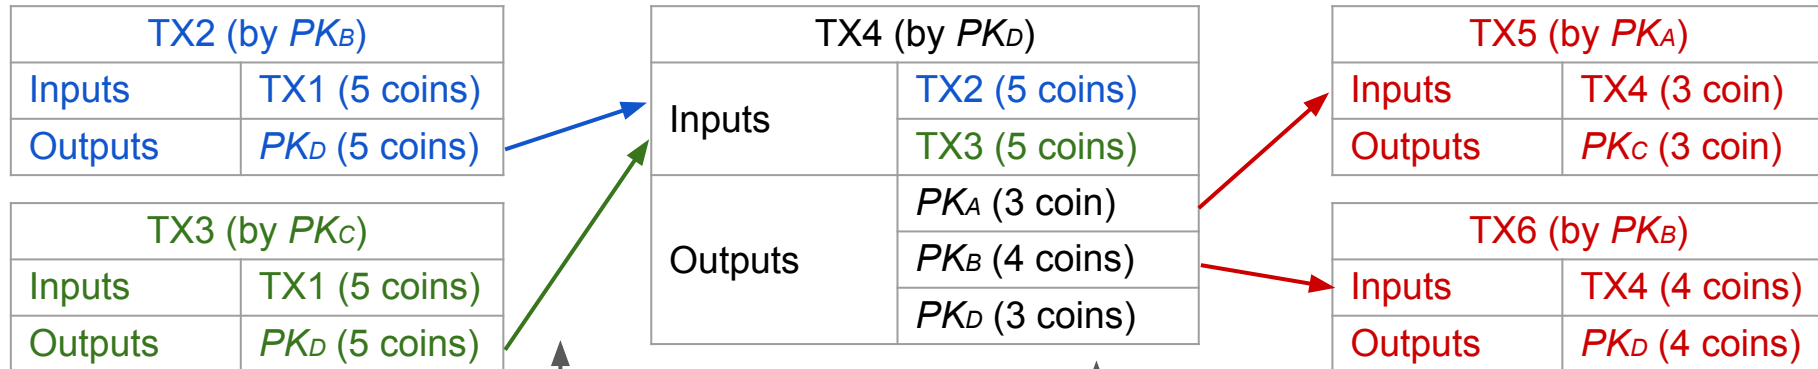
Transactions



The Ledger v2

1. PK_B and PK_C magically start with 5 coins.
2. $\{\text{"Using currency from TX 1, give 5 coins to } PK_D.\text{"}\}_{SK_B^{-1}}$
3. $\{\text{"Using currency from TX 1, give 5 coins to } PK_D.\text{"}\}_{SK_C^{-1}}$
4. $\{\text{"Using currency from TX 2 and TX 3, give 3 coin to } PK_A \text{ and 4 coins to } PK_B, \text{ and 3 coins to } PK_D\text{"}\}_{SK_D^{-1}}$
5. $\{\text{"Using currency from TX 4, give 3 coin to } PK_C.\text{"}\}_{SK_A^{-1}}$
6. $\{\text{"Using currency from TX 4, give 4 coins to } PK_D.\text{"}\}_{SK_B^{-1}}$

Transactions



Each transaction must identify a valid source of currency by referencing a past transaction.


We must also verify that no other transaction has already claimed this input.

In each transaction, the total input currency is equal* to the total output currency.

In TX 4, $5 + 5 = 3 + 4 + 3$

Leftover money is returned to the sender. PK_D spent 3 coins and sent the remaining 7 back to himself.

Bitcoin: The Public Ledger




Textbook Chapter 16.6–16.9

Recall: Hash Functions

- Hash functions produce a fixed-length “fingerprint” over an arbitrary length of data
 - **Preimage resistant:** Given an output, difficult to find an input that hashes to the output
 - **Collision resistant:** Difficult to find two inputs that hash to the same output
- In practice, hash functions “look” random
 - Changing the input causes the output to change unpredictably
 - Each bit in the output has a 50% chance of flipping

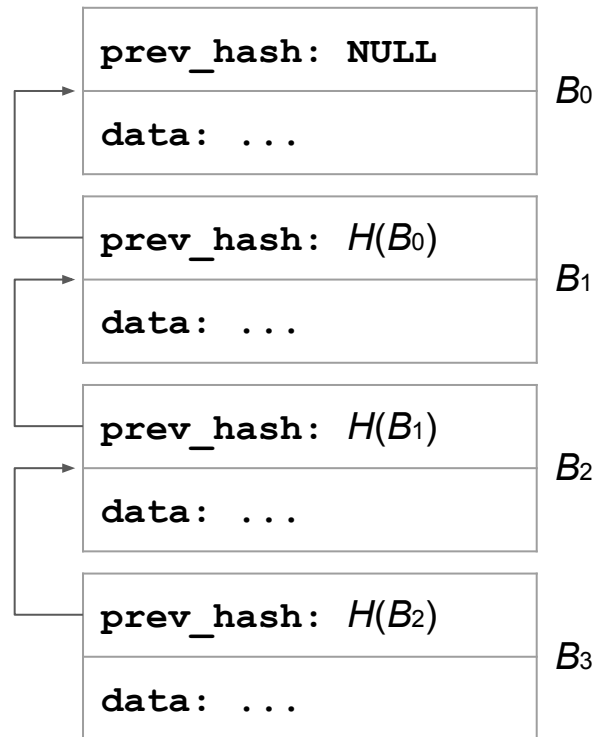
Hash Chains

- What we want: A data structure where we can append a node and then compute a hash over all nodes efficiently
 - Appending a node of size n should take $O(n)$ time
- Idea: To validate the previous block, include a hash of the previous block
 - The previous block validates the block before, etc.
- To append:
 - Compute the hash of the current block
 - Construct a new block containing the previous hash and the data
 - Set the head of the chain to the new block



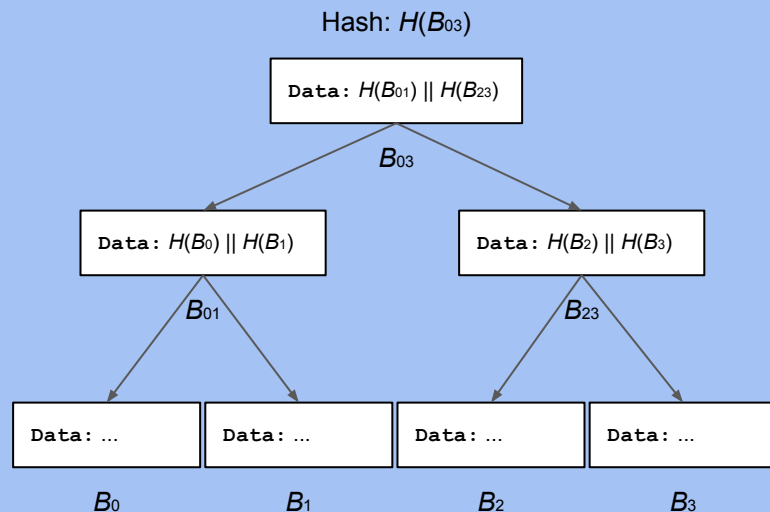
Hash Chains

- Result
 - The latest hash represents a hash over all previous nodes
 - Changing data changes the block's hash, which changes the next block's hash, etc.
- This is really just an append-only linked list
 - Git uses this: Each commit contains a hash of the previous commit



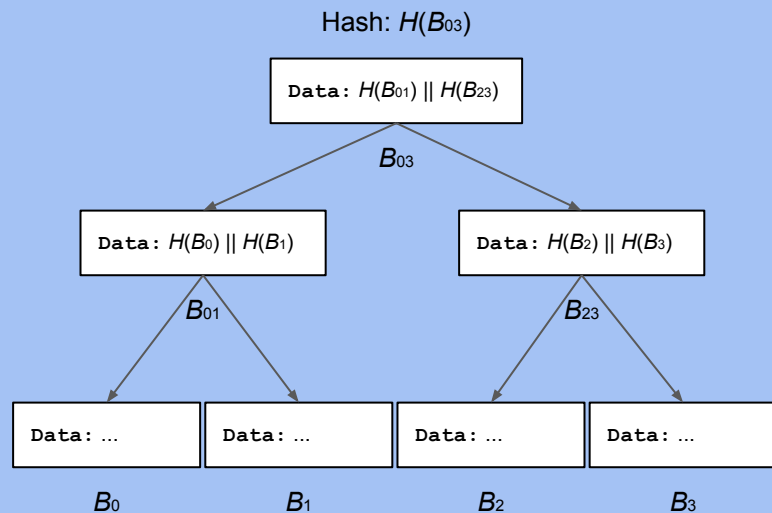
Merkle Trees

- What we want: A data structure where we can modify a node and then compute a hash over all nodes efficiently
 - Appending a node of size n to a structure with m nodes should take $O(n + \log m)$ time
- Idea: Instead of hashing all nodes at once, combine hashes as a binary tree
 - Each node is a hash of the two child node's hashes
- To modify:
 - Modify and re-hash the block
 - Re-hash the parent nodes until you reach the root



Merkle Trees

- Result
 - The top hash represents a hash over all nodes
 - Can easily modify any node and compute $O(\log m)$ hashes in order to compute a new top hash
- Bitcoin uses a linear chain instead of a tree, but Merkle trees appear in many other applications




The Public Ledger

- Our previous scheme needed a public, trusted ledger
 - Append-only, immutable
 - Accessible to all parties
 - We can't trust any individual to maintain the ledger for us
- **Consensus:** A way to reach agreement on something as a group, without placing trust in any individual
 - For simplicity, combine multiple transactions into one block, so we only need to reach consensus for one block at a time while maintaining several hundred transactions

Forking Attacks

- Do hash chains allow a group of individuals to reach a consensus?
 - **No.** If Mallory creates an alternate history where she didn't spend 100,000 coins, this is still a valid hash chain!
 - Need a way to agree on the head of the blockchain



Proof of Work

- In each block, include a random number (“nonce”)
- For a block to be valid, the hash of the block must start with some number of 0’s
 - If you want 4 0’s, each nonce has a $1/2^4 = 1/16$ chance of creating a valid block
- Each user keeps track of their own blockchain

Hash: 0b11100011

```
nonce: 0x0000
prev_hash: ...
data: ...
```



Hash: 0b01001000

```
nonce: 0x0001
prev_hash: ...
data: ...
```



Hash: 0b00011010

```
nonce: 0x0002
prev_hash: ...
data: ...
```



Hash: 0b00000101

```
nonce: 0x0017
prev_hash: ...
data: ...
```



Proof of Work

- Users “mine” a block by trying nonce values until a valid block is found
 - This requires trying many nonces to find one to produce the correct hash value for the block
 - The valid block is broadcasted to everyone else on the network
- When a valid block is received, users add it to their blockchain and start mining the next block

Hash: 0b11100011

```
nonce: 0x0000
prev_hash: ...
data: ...
```



Hash: 0b01001000

```
nonce: 0x0001
prev_hash: ...
data: ...
```



Hash: 0b00011010

```
nonce: 0x0002
prev_hash: ...
data: ...
```



Hash: 0b00000101

```
nonce: 0x0017
prev_hash: ...
data: ...
```



Proof of Work

- **Consensus: The longest blockchain is the “true” blockchain**
 - Everyone mines on the longest chain they know of
 - If we hear about a longer chain, we discard our current mining block and begin mining the longer chain
 - The more people agree to a chain, the more mining is done on that chain (more hashes per second)
 - The more mining is done on a chain, the faster the chain grows
- **Foiling the forking attack**
 - For Mallory to fork the blockchain, she must mine her new chain faster than the current “true” chain
 - To mine a chain faster than the current “true” chain, she must mine faster than every other (honest) node combined
 - This is infeasible!

Proof of Work: Examples

- What if two miners mine two blocks at the same time and append it to the blockchain, causing a fork?
 - The next miner that appends onto one of these chains invalidates the other chain. Longest chain wins.
- If a miner included your transaction in the latest block created, are you guaranteed that your transaction is forever in the blockchain?
 - No. There could have been another miner appending a different block at the same time, and that chain might be winning.
 - To confirm a transaction, wait for a couple more blocks to be appended to the chain with your transaction to ensure that it will probably win in the long term
- What happens if a miner who just mined a block refuses to include my transaction?
 - Hopefully, the next miner to mine a block is willing to include your transaction.

Incentivizing Mining

- Miners are incentivized to mine blocks in exchange for currency
 - Each signed transaction includes a small fee given to the miner of a block
 - Example: "... and pay 0.0001 coins to the miner of this transaction"
 - Technically, the transaction fee paid is (sum of inputs) - (sum of outputs)
 - Each block may also give an agreed-upon number of free coins for the miner of the block
 - Example: Block B has 3 transactions and an additional transaction that reads, " PK_A receives 25 free coins"
- As more miners join the pool, the algorithm adjusts the number of 0's needed to mine a block
 - Generally, the time per block is targeted at a certain amount of time, so more global hash power \Rightarrow more 0's to make mining harder

Attacks on Proof-of-Work

- **51% attack:** An attacker who controls 51% of mining power can effectively rewrite history and perform a forking attack
 - Creates a centralized authority: An entity is given control over the currency network
 - Proof-of-work makes this difficult: The attacker must control 51% of the world's computing power
- Bitcoin's security relies on the assumption that no attacker controls 51% of the world's computing power




The Trouble with Bitcoin



The Trouble with Bitcoin: Centralization of Power

- Bitcoin is often not decentralized in practice
 - Recall: In Bitcoin's design, there is no centralized authority controlling the system
 - In practice, a few small groups have a lot of control over the Bitcoin system
- **Mining pools:** A team of users mining together
 - When one user receives a mining reward, everyone in the team shares the reward together
 - A user mining alone must get lucky to receive a reward
 - A mining pool gives users steady, smaller rewards
- A few large mining pools control most of the computing power in Bitcoin
 - If large pools team up, the 51% attack is possible!



The four largest mining pools combined control 55% of all Bitcoin hash power

The Trouble with Bitcoin: Centralization of Power

- Codebases are developed and maintained by a few groups
 - These groups can rewrite the code to affect the entire system
 - Example: When Ethereum was hacked, the developers changed the code to retrieve their money
- Some cryptocurrencies are not decentralized by design
 - **Private blockchain:** Only blocks signed by trusted private keys are valid
 - Only central authorities can sign blocks, but anyone can validate
 - Defeats the point of decentralized consensus!

The Trouble with Bitcoin: Pseudonymity


- Bitcoin is *pseudonymous*
 - **Pseudonymity:** Multiple actions can be linked to a single identity which is not your real identity, a pseudonym
 - The pseudonym is the public key
 - Your transactions can be linked to your public key
- Bitcoin transactions are not necessarily *anonymous*
 - **Anonymity:** Actions cannot be linked to your real identity
 - In theory, if you only ever use Bitcoin in an unpredictable manner, your pseudonym cannot be linked to your identity (anonymous)
 - In practice, your pseudonym can be linked to your real identity
 - If your shopping habits are predictable, it will be linked to you (the ledger is public)
 - If you exchange Bitcoin with real currency (dollars, euros, etc.), it will be linked to you

The Trouble with Bitcoin: Inefficiency

- Runtime inefficiency
 - Proof-of-work requires a huge amount of useless computational work
 - Bitcoin: Hashing random values until a hash starts with many zeros
- Storage inefficiency
 - Each Bitcoin user must store the entire transaction history to validate transactions
 - Alternative: Don't store the entire transaction history
 - Problem: Now you have to ask a trusted source to send you the history
 - Defeats the point of decentralization!
- Result: Low processing capacity
 - Original design limited each block to 1 MB in size to defend against spam
 - Bitcoin can only process 3–7 transactions per second

The Trouble with Bitcoin: Power Consumption

- Bitcoin requires computers to waste energy on proof-of-work computations
- Bitcoin mining incentivizes energy waste
 - People seek mining rewards
 - If mining becomes more efficient, Bitcoin adjusts the mining problem to be harder, so efficiency gains are not realized



Today, Bitcoin consumes more electricity than entire countries (measured in terawatt-hours)

The Trouble with Bitcoin: Irreversibility

- Modern banking transactions are designed to be reversible
 - Protects against fraud: If your money is stolen, the bank can recover the money for you
- Bitcoin is not reversible
 - Once Bitcoin is sent, the transaction cannot be undone (unless the sender wants to return the money to you)
- You only need your private key to access your Bitcoin
 - If someone steals your private key, they can access all your money
 - If your Bitcoin is stolen, there is no way to recover it
 - Result: It is dangerous to store Bitcoin on any computer connected to the Internet



Bitcoin in Practice

Breaking Proof-of-Work

- Proof-of-work viewed from an economic perspective
 - The system wastes $\$x$ per hour to defend against attacks
 - Each honest user contributes a small amount of the $\$x$
 - Inefficiency: Money is constantly wasted, even if the service is not under attack
 - An attacker must spend more than $\$x$ per hour to control the system
 - Assumption: An attacker will not spend more than $\$x$ per hour
 - If an attacker can make more than $\$x$ per hour for an attack, they will spend $\$x$ per hour to execute the attack!
 - There are hashing services the attacker can pay to execute the attack

Proof-of-Work Alternatives

- Main problem: Proof-of-work is inefficient and wastes electricity
 - Proof-of-work is necessary to prevent an attacker controlling the system (makes it expensive to control the system)
- Alternative: Proof-of-stake
 - Proof-of-stake: Users with more coins have more mining power
 - People coins *stake* them in the system
 - If they act honestly, they gain more coins
 - If they act maliciously, their stake is slashed
 - Benefit: Uses much less electricity than proof-of-work, and the attacker must burn their own coins to attack the system!
 - Problem: Gives more power to wealthier users
- Alternative: Articulated trust
 - Articulated trust: Designate several trusted parties, who vote on each transaction
 - As long as half of the trusted parties are trustworthy, then the protocol is secure
 - Relying on a little bit of trust solves both efficiency and energy problems!

Blockchain: Marketing and Buzzwords

- “Blockchain” is often marketed as brand-new technology, but it is mostly existing technology
 - Hash chains are over 20 years old: Linked timestamping services used hash chains and were proposed in 1990
 - Merkle trees were patented in 1979
 - Private blockchains are not new technology: Many existing applications already use append-only database structures
- “Blockchain” is a buzzword often applied to completely unrelated problems
 - “Use blockchain for electronic voting”
 - “Use blockchain to store medical records”
 - “Use blockchain to deliver vaccines”

Blockchain: Marketing and Buzzwords

- Example of blockchain marketing: “Smart Contracts”
 - Smart contracts: Write an agreement in code, and execute the code to automatically follow the procedure
 - The code is stored on a blockchain so it cannot be modified
 - Problems with smart contracts
 - Unclear if they actually solve any real problem
 - Most “smart contracts” are actually standard finance bots: small programs to perform money transfers
 - If the code is vulnerable, you can violate the contract, with no way to reverse the effects
 - Contrast with legal contracts: Issues are handled by court systems

Blockchain: Marketing and Buzzwords

- Example of blockchain marketing: Non-Fungible Tokens (NFTs)
- NFT: A piece of data certifying that a digital file belongs to someone
 - Basically the statement $\{\text{"This work of art now belongs to \$individual"}\}_{PK_{\text{prev_owner}}}$ stored on the blockchain
 - Unrelated to copyrights or digital sharing of the file
- The NFT market became extremely valuable in 2020–2021
 - NFTs of digital files have been selling for millions of dollars
 - Blockchain marketing: Adding a blockchain didn't solve anything, but the "new technology" convinced buyers
 - Speculative bubble: NFTs are not worth anything on their own, but people buy NFTs to sell them later for more money, resulting in price increases

Bitcoin Enables Censorship Resistance

- Bitcoin has no central authority to block transactions
 - Bitcoin can be used for electronic payments that standard platforms would block
- Wikileaks: An organization that publishes leaked classified information
 - Used by whistleblowers to expose government corruption and corporate wrongdoing
 - Opposed and censored by many governments
- Bitcoin is used to support Wikileaks
 - Many major platforms refused any donations to Wikileaks
 - Bitcoin has no central authority, so nobody can stop the Wikileaks donations
- This is, generally, a good thing! But...

Bitcoin Enables Crime

- Bitcoin is used for illegal transactions
 - Drug dealing
 - Money laundering
 - Illegal gambling
 - Hiring hitmen
 - Ransomware and extortion
- Bitcoin has no central authority to block illegal transactions
- Bitcoin is the most effective way to make illegal transactions

Economics of Bitcoin: Volatility

- Volatile currency: The value changes quickly
 - The value of Bitcoin changes far more often than standard currency
- Bitcoin is vulnerable to price shocks
 - Price shock: An extremely sudden change in value
 - When there are more transactions than the block capacity, prices increase
 - Users are competing for a limited number of transactions
 - Unknown attacks have also caused price shocks
- Result: Bitcoin behaves more like stock than currency
 - Users keep Bitcoin to try and grow their investment when the value of Bitcoin increases

Economics of Bitcoin: Speculation

- Because it is so volatile, Bitcoin behaves more like stock than currency
- Speculation: Buying something so that you can sell it later for more money
 - You don't buy Bitcoin because owning Bitcoin helps you make money
 - You buy Bitcoin because you hope to sell it later for more money
 - Relies on short-term price changes and not long-term value
- Speculation results in a bubble
 - Bubble: Something sells for more than its true value
 - As more people buy Bitcoin to try and make a profit, the price of Bitcoin also increases
 - Bubbles always burst: Eventually the price returns to its original value, leading to huge economic losses

Economics of Bitcoin: Currency Exchange

- Companies and people prefer to keep money in a more stable currency
 - To buy a product in Bitcoin, the buyer converts their standard currency to Bitcoin
 - The seller receives the Bitcoin and immediately converts it back to standard currency
- Users should be able to exchange Bitcoin for other currency (e.g. dollars)
- Buying and selling Bitcoin is difficult
 - Recall: Bitcoin transactions are irreversible
 - The buyer must trust that the seller will transfer the Bitcoin
 - The seller must trust that the buyer will pay when the Bitcoin is transferred
 - Ways to buy Bitcoin
 - Use another irreversible payment (e.g. cash)
 - Have a trusted relationship with the seller
 - Send a deposit first

Economics of Bitcoin: Volatility



- Stable currencies require reliable conversion to other currency
 - Someone must support easy conversion between Bitcoin and other currency
 - This is usually a centralized entity (e.g. a bank or a government)
 - The centralized entity violates Bitcoin's main purpose
- Options for operating a conversion service
 - Follow government regulations (essentially becoming a regular currency)
 - Operate independently as a “wildcat bank” (similar to banks in the 1800s, before the US had a national currency)
 - Ignore federal regulations (e.g. Liberty Reserve, which was shut down by the government)
- Some cryptocurrency designs claim to be “algorithmic stablecoins”
 - Designed to be stable in the market
 - Most of these are snake oil that don't work

Other Cryptocurrencies

- Most cryptocurrencies are based on the same principles
 - Public, append-only ledger structure
 - Designed with decentralization in mind
 - Some are software “forks” of the original Bitcoin blockchain
 - The fork ignores new Bitcoin blocks, and Bitcoin ignores fork blocks
- New cryptocurrencies are marketed with a distinguishing feature
 - Litecoin: Adds a catchy slogan
 - Dogecoin: Adds an Internet meme
 - Ripple: Centralized cryptocurrency with an additional settlement structure
 - IOTA: Designed its own brand-new cryptography (using trinary math)
 - Monero: Improves pseudonymity
 - Zcash: Adds real anonymity
 - Ethereum: Adds million-dollar rewards for catching bugs

Other Cryptocurrencies

- Most cryptocurrencies are based on the same principles
 - Public, append-only ledger structure
 - Designed with decentralization in mind
 - Some are software “forks” of the original Bitcoin blockchain
 - The fork ignores new Bitcoin blocks, and Bitcoin ignores fork blocks
- New cryptocurrencies are marketed with a distinguishing feature
 - Litecoin: Adds a catchy slogan
 - Dogecoin: Adds an Internet meme
 - Ripple: Centralized cryptocurrency with an additional settlement structure
 - IOTA: Designed its own brand-new cryptography (using trinary math)
 - Monero: Improves pseudonymity
 - Zcash: Adds real anonymity
 - Ethereum: Adds million-dollar rewards for catching bugs



Cryptocurrency Scams

- Public interest in cryptocurrency as a “get rich quick” scheme leads to fraud
- Many cryptocurrency frauds are old frauds with new technological branding
 - Ponzi schemes: Trick uninformed consumers to invest in a nonexistent product
 - Some “smart contracts” are actually modern Ponzi schemes
 - Wildcat banks: Independent, unregulated banks
 - If the bank shuts down, your money in the bank is gone
 - Physical wildcat banks stopped existing in the 1800s, but some cryptocurrencies essentially operate as wildcat banks
 - Unregulated securities: Stocks that are not regulated by the government
 - Often scams: The unregulated stock may be completely worthless
 - Initial coin offerings: Pay money now in exchange for some coins when the cryptocurrency launches later
 - If the cryptocurrency never launches, your money is gone

Bitcoin: Summary

- Goal: Create a currency system that does not rely on any central authority
- Identity: Each user is identified by their public key
- Transactions
 - Users sign transactions with their private key and add them to the ledger
 - Each transaction must reference a previous transaction to identify a source of money
- Public ledger
 - Hash chain: A linked list where each node contains the hash of the previous node
 - Append-only structure: Changing a node causes the hashes in all future nodes to change
 - Vulnerable to forking attacks: The attacker creates their own branch of the chain
- Proof-of-work
 - The blockchain only accepts blocks whose hash starts with a sequence of n 0s
 - Finding valid blocks requires trying 2^n hashes. A reward is given to incentivize mining blocks
 - The longest hash chain is accepted as the true blockchain
 - An attacker must control 51% of the world's computing power to create their own hash chain

The Trouble with Bitcoin: Summary

- **Centralization of power:** In practice, Bitcoin is controlled by a few groups
 - Mining pools: Teams of users mining blocks together
 - Codebase developers: Can change the code to alter the system
 - Private blockchains: Only trusted parties can append to the blockchain
- **Pseudonymity**
 - In theory, your transactions are only linked to your public key, not your true identity
 - With predictable transactions, your public key can be linked to your identity too
- **Inefficiency**
 - Proof-of-work requires a huge amount of hashing
 - Each user must store the entire blockchain
 - Bitcoin can only process a few transactions per second
- **Power consumption:** Hashing wastes electricity
- **Irreversibility:** Transactions are not reversible
 - If your Bitcoin is stolen, there is no way to recover it