

Transport Layer: TCP and UDP

CS 161 Fall 2022 - Lecture 18

Last Time: Low-Level Network Attacks

Computer Science 161

- **Classes of attackers:**
 - Off-path: Can't see, modify, or drop packets
 - On-path: Can see packets, but can't modify or drop packets
 - MITM: Can see, modify, and drop packets
- **ARP: A protocol to translate local IP addresses to MAC addresses**
 - Ask everyone on the network, "Who has the IP 1.2.3.4?"
 - Attack: The attacker can respond instead of the true device with 1.2.3.4, and packets will get routed to the attacker!
 - Defense: Switches
 - Defense: Rely on higher layers
- **DHCP: A protocol for a new client to receive a network configuration**
 - Ask everyone on the network, "What is the network configuration to use?"
 - Attack: The attacker can respond with a malicious configuration
 - Defense: Rely on higher layers

Last Time: Low-Level Network Attacks

- WPA: A protocol to encrypt Wi-Fi connections at layer 1
 - Messages between the client and the AP are encrypted with keys
 - Handshake uses MICs (cryptographic MACs) to verify that both parties have the same PSK and nonces
 - WPA-PSK: Use a password to derive a PSK, which is used in a handshake to arrive at a key
 - Attack: Attacker can pretend to be an AP
 - Attack: Brute-force the password after recording a handshake
 - Vulnerability: No forward secrecy
 - WPA-Enterprise: Use a third party to provide a one-time “replacement PSK,” used in the same handshake
 - Solves the attacks on WPA-PSK

Last Time: BGP

- Border Gateway Protocol (BGP): Routing packets
 - The Internet is made of smaller **autonomous systems (AS)**
 - Each AS broadcasts the shortest routes it knows of (dependent on the shortest routes of its neighbors and distance to neighbors)

Today: Transport Layer Protocols

Computer Science 161

- Transmission Control Protocol (TCP): Reliably sending packets
- User Datagram Protocol (UDP): Non-reliably sending packets

Transmission Control Protocol (TCP)

Textbook Chapter 30

Review: IP Reliability

- **Reliability** ensures that packets are received correctly
 - If the packet has been changed (random error or malicious tampering), it should not be correctly received
 - IP packets include a checksum (unkeyed function, protects against random errors)
 - However, there is no cryptographic MAC, so there are no guarantees if an attacker maliciously modifies packets (and modifies the checksum accordingly)
- IP is **unreliable** and only provides a **best effort** delivery service, which means:
 - Packets may be lost (“dropped”)
 - Packets may be corrupted
 - Packets may be delivered out of order
- It is up to higher level protocols to ensure that the connection is reliable

Scratchpad: Let's Design It Together

- Problem: IP packets have a limited size. To send longer messages, we have to manually break messages into packets
 - When sending packets: TCP will automatically split up messages
 - When receiving packets: TCP will automatically reassemble the packets
 - Now the user doesn't need to manually split up messages!
- Problem: Packets can arrive out of order
 - When sending packets: TCP labels each byte of the message with increasing numbers
 - When receiving packets: TCP can use the numbers to rearrange bytes in the correct order
- Problem: Packets can be dropped
 - When receiving packets: TCP sends an extra message acknowledging that a packet has been received
 - When sending packets: If the acknowledgement doesn't arrive, re-send the packet

Transmission Control Protocol (TCP)

- Provides a byte stream abstraction
 - Bytes go in one end of the stream at the source and come out at the other end at the destination
 - TCP automatically breaks streams into **segments**, which are sent as layer 3 packets
- Provides ordering
 - Segments contain sequence numbers, so the destination can reassemble the stream in order
- Provides reliability
 - The destination sends **acknowledgements** (ACKs) for each sequence number received
 - If the source doesn't receive the ACK, the source sends the packet again
- Provides ports
 - Multiple services can share the same IP address by using different ports

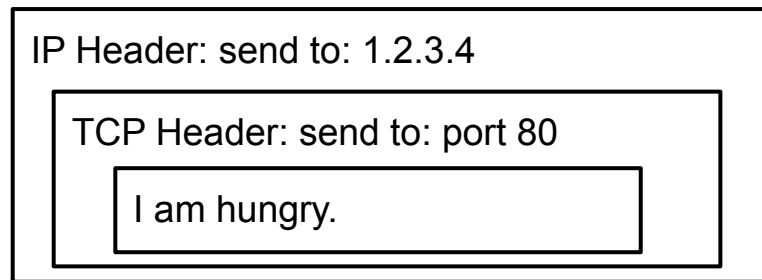
Ports: An Analogy

Computer Science 161

- Alice is pen pals with Bob. Alice's roommate, Carol, is also pen pals with Bob
- Bob's replies are addressed to the same global (IP) address
 - How can we tell which letters are for Alice and which are for Bob?
- Solution: Add a room number (port number) inside the letter
 - In private homes, usually a port number is meaningless
 - But, in public offices (servers), like Cory Hall, the port numbers are constant and known

Ports

- **Ports** help us distinguish between different applications on the same computer or server
 - On private computers, port numbers can be random
 - On public servers, port numbers should be constant and well-known (so users can access the right port)
- Remember: TCP is built on top of IP, so the IP header (and therefore the IP address) is still present



Establishing Sequence Numbers

- Each TCP connection requires two sets of sequence numbers
 - One sequence number for messages from the client to the server
 - One sequence number for messages from the server to the client
- Before starting a TCP connection, the client and server must agree on two **initial sequence numbers (ISNs)**
 - The ISNs are different and random for every connection (for security reasons, as we'll see soon)

H	e	l	l	o		s	e	r	v	e	r
50	51	52	53	54	55	56	57	58	59	60	61


Messages from the client are numbered starting at 50.

H	e	l	l	o		c	l	i	e	n	t
25	26	27	28	29	30	31	32	33	34	35	36

Messages from the server are numbered starting at 25.

TCP: 3-Way Handshake

1. Client chooses an initial sequence number x its bytes and sends a SYN (synchronize) packet to the server
2. Server chooses an initial sequence number y for its bytes and responds with a SYN-ACK packet
3. Client then returns with an ACK packet
4. Once both hosts have synchronized sequence numbers, the connection is “established”



TCP: Sending and Receiving Data

- The TCP handlers on each side track which TCP segments have been received for each connection
 - A connection is identified by these 5 values (sometimes called a 5-tuple)
 - Source IP
 - Destination IP
 - Source Port
 - Destination Port
 - Protocol
- Data from the bytestream can be presented to the application when all data before has been received and presented
 - Recall: TCP presents data to the application as a bytestream, so the order must be preserved from one end to the other, even if packets are received out of order

TCP: Sending and Receiving Data

- Byte i of the bytestream is represented by sequence number $x + i$
 - The first byte is byte $i = 1$, since sequence number x was used for the SYN packet and y for the SYN-ACK packet
- A packet's sequence number is the number of the first byte of its data
 - This number is from the sender's set of sequence numbers
- A packet's ACK number, if the ACK flag is set, is the number of the byte immediately after the last received byte
 - This number is from the receiver's set of sequence numbers
 - This would be (sequence number) + (length of data) for the last received packet

TCP: Sending and Receiving Data



TCP: Retransmission

- If a packet is dropped (lost in transit):
 - The recipient will not send an ACK, so the sender will not receive the ACK
 - The sender repeatedly tries to send the packet again until it receives the ACK
- If a packet is received, but the ACK is dropped:
 - The sender tries to send the packet again since it didn't receive the ACK
 - The recipient ignores the duplicate data and sends the ACK again
- When packets are dropped in TCP, TCP assumes that there is congestion and sends the data at a slower rate

TCP: Ending/Aborting a Connection

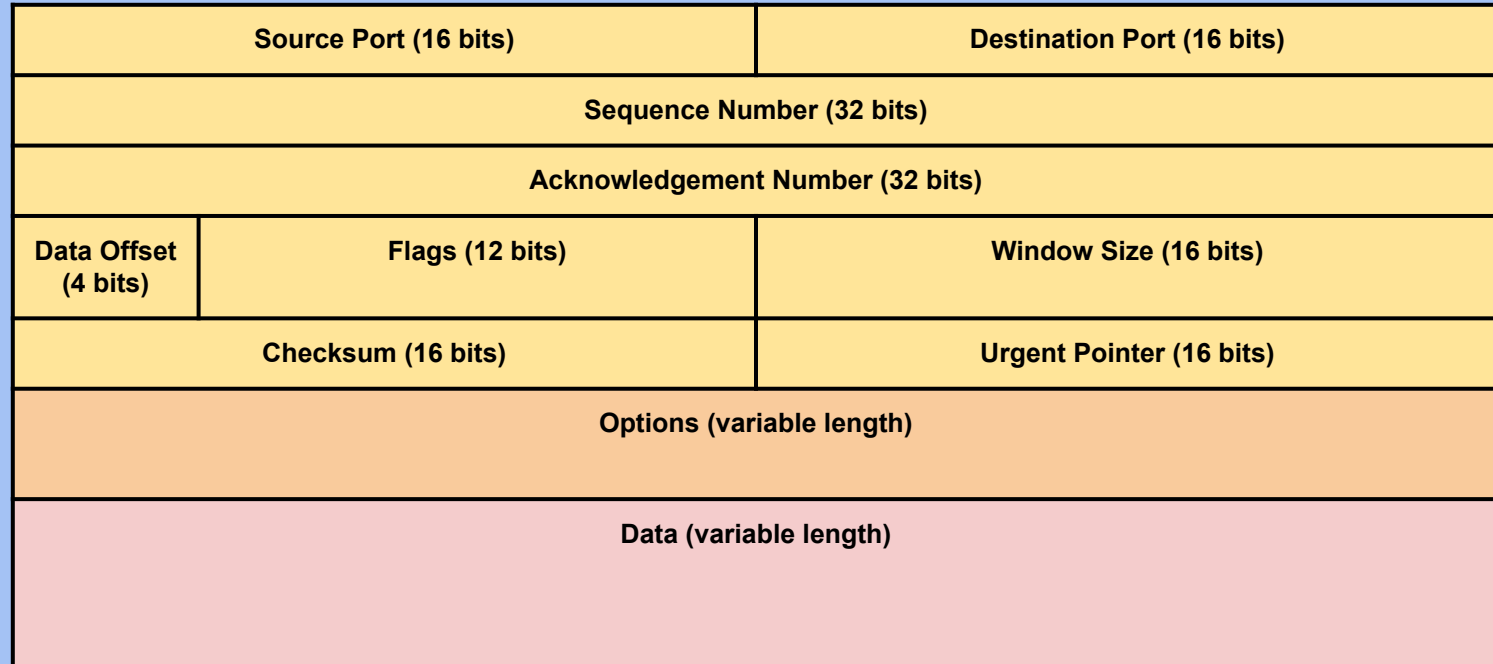
- To **end** a connection, one side sends a packet with the FIN (finish) flag set, which should then be acknowledged
 - This means “I will no longer be sending any more packets, but I will continue to receive packets”
 - Once the other side is no longer sending packets, it sends a packet with the FIN flag set
- To **abort** a connection, one side sends a packet with the RST (reset) flag set
 - This means “I will no longer be sending nor receiving packets on this connection”
 - RST packets are not acknowledged since they usually mean that something went wrong

TCP Flags

- **ACK**
 - Indicator that the user is acknowledging the receipt of something (in the ack number)
 - Pretty much always set except the very first packet
- **SYN**
 - Indicator of the beginning of the connection
- **FIN**
 - One way to end the connection
 - Requires an acknowledgement
 - No longer sending packets, but will continue to receive
- **RST**
 - One way to end a connection
 - Does not require an acknowledgement
 - No longer sending or receiving packets

TCP Packet Structure

Computer Science 161




TCP segment header

TCP Attacks

- **TCP hijacking:** Tampering with an existing session to modify or inject data into a connection
 - **Data injection:** Spoofing packets to inject malicious data into a connection
 - Need to know: The sender's sequence number
 - Easy for MITM and on-path attackers, but off-path attackers must guess 32-bit sequence number (called **blind injection/hijacking**, considered difficult)
 - For on-path attackers, this becomes a race condition since they must beat the server's legitimate response
 - **RST injection:** Spoofing a RST packet to forcibly terminate a connection
 - Same requirements as packet injection, so easy for on-path and MITM attackers, but hard for off-path attackers
 - Often used in censorship scenarios to block access to sites

TCP Data Injection

Computer Science 161

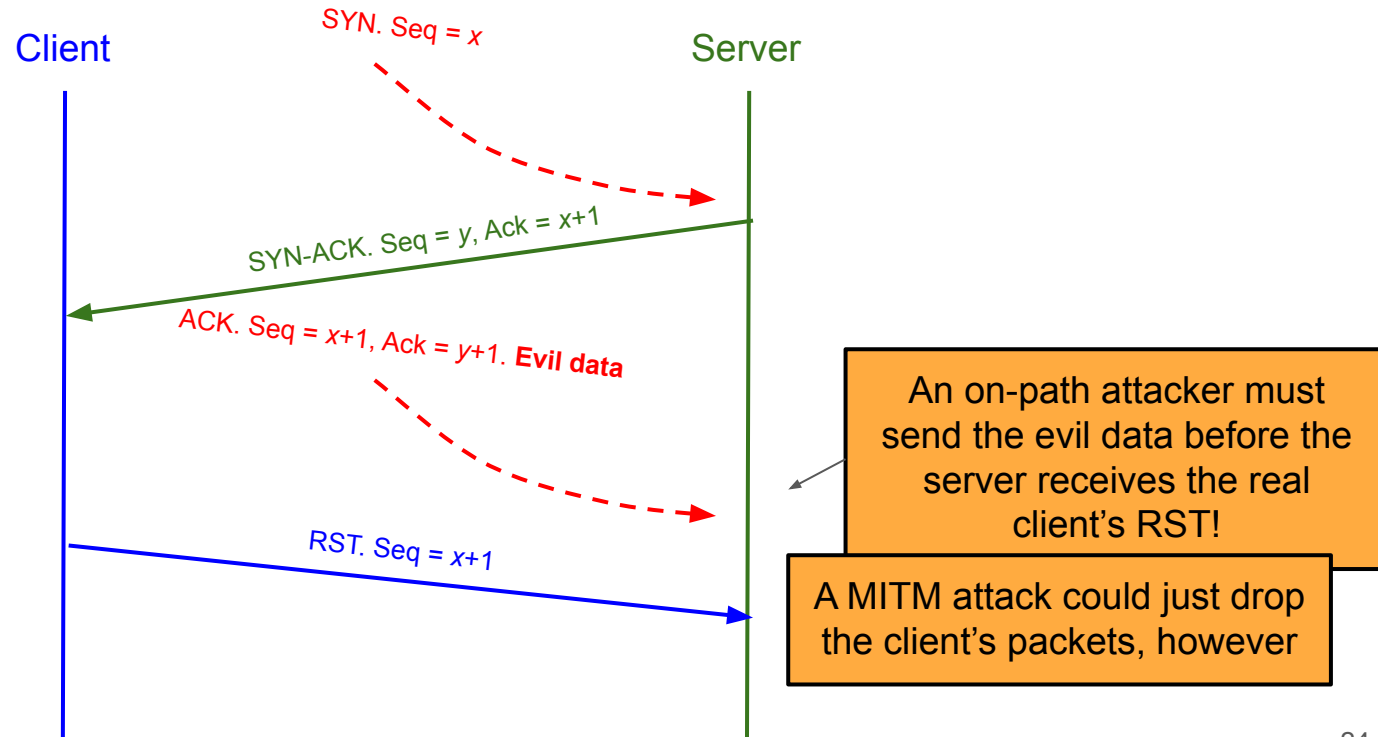


This packet will be ignored by the client since the client already processed the malicious packet!

TCP Attacks

- **TCP spoofing:** Spoofing a TCP connection to appear to come from another source IP address
 - Recall: IP packets can often be spoofed if the AS doesn't block source addresses
 - Need to know: Sequence number in the server's response SYN-ACK packet
 - Easy for MITM and on-path attackers, but off-path attackers must guess 32-bit sequence number (called **blind spoofing**, also considered difficult)
 - For on-path attackers, this is a race condition, since the real client will send a RST upon receiving the server's SYN-ACK!

TCP Spoofing



TCP Attacks

- TCP provides no confidentiality or integrity
 - Instead, we rely on higher layers (like TLS, more on this next time) to prevent those kind of attacks
- Defense against off-path attackers rely on choosing random sequence numbers
 - Bad randomness can lead to trivial off-path attacks: TCP sequence numbers used to be based on the system clock!

User Datagram Protocol (UDP)

Textbook Chapter 30

User Datagram Protocol (UDP)

- Provides a **datagram** abstraction
 - A message, sent in a single layer 3 packet (though layer 3 could fragment the packet)
 - Max size limited by max size of packet
 - Applications break their data into datagrams, which are sent and received as a single unit
 - Contrast with TCP, where the application can use a bytestream abstraction
- No reliability or ordering guarantees, but adds ports
 - It still has *best effort* delivery
- Much faster than TCP, since there is no 3-way handshake
 - Usually used by low-latency, high-speed applications where errors are okay (e.g. video streaming, games)

UDP Attacks

- No sequence numbers, so relatively easy to inject data into a connection or spoof connections
 - Higher layers must provide their own defenses against these attacks!

UDP Packet Structure

Computer Science 161

Source Port (16 bits)	Destination Port (16 bits)
Length (16 bits)	Checksum (16 bits)
Data (variable length)	

UDP datagram header

Summary

- **Transmission Control Protocol (TCP): Reliably sending packets**
 - 3-way handshake: Client sends SYN, server sends SYN-ACK, client sends ACK
 - Provides reliability, ordering, and ports
 - Attack: TCP hijacking through data injection or RST injection
 - Blind attacks must guess the client's or server's sequence numbers
 - Attack: TCP spoofing by sending a spoofed SYN packet
 - Blind attacks must guess the server's sequence number
- **User Datagram Protocol (UDP): Non-reliably sending packets**
 - No reliability or ordering, only ports
 - Same injection and spoofing attacks as TCP, but easier